

# Predictive Model


## Overview

The **Predictive Policing System** is a data-driven tool designed to assist command staff and dispatchers in allocating patrol resources more effectively. By analyzing historical crime data, the system predicts high-risk zones for violent crime in 1-hour windows, allowing for proactive rather than reactive policing.

## How to Use the System

### A. Accessing the Dashboard

1. Open your web browser and navigate to the local server address:  
`http://127.0.0.1:8000/docs`.
2. You will see the **Swagger UI Dashboard**, which serves as the control panel for the system.



### B. Checking System Status

1. Click on the blue bar labeled `GET /`.
2. Click **Try it out → Execute**.
3. The system will return a "200 OK" response, confirming that the machine learning model is loaded and ready.

## C. Generating a Risk Prediction

To check the risk level for a specific district and time:

1. Click the green bar labeled **POST /predict**.
2. Click the **Try it out** button on the right.
3. Enter the shift details in the JSON box:
  - `district` : The Police District ID (e.g., `11`).
  - `hour` : The hour of the shift in 24-hour format (e.g., `23` for 11 PM).
  - `day_of_week` : The day (0=Monday, ... 5=Saturday, 6=Sunday).
  - `month` : The month number (e.g., `12` for December).
4. Click the blue **Execute** button.

The screenshot shows a UI for generating a risk prediction. At the top, there's a green bar with the text "Parameters" and a red "POST /predict" button. Below this is a "Request body" section with a "required" label and a dropdown set to "application/json". The main area contains a JSON editor with the following content:

```
{  
  "district": 11,  
  "hour": 23,  
  "day_of_week": 5,  
  "month": 12  
}
```

At the bottom is a blue "Execute" button.

## D. Interpreting the Results

Scroll down to the **Server response** section. You will see a result like this:

JSON

```
{  
  "district": 11,  
  "risk_score": 0.7421,  
  "alert": true,  
  "action": "DISPATCH"  
}
```

- **Risk Score:** The calculated probability (0-1) of violent crime occurring.

- **Alert (True/False):** Indicates if the risk exceeds the safety threshold.
- **Action:**
  - **DISPATCH:** High risk detected. Recommend deploying a tactical unit.
  - **MONITOR:** Standard risk. Routine patrol is sufficient.

```

Curl
curl -X 'POST' \
  'http://127.0.0.1:8000/predict' \
  -H 'accept: application/json' \
  -H 'Content-type: application/json' \
  -d '{
    "district": 11,
    "hour": 23,
    "day_of_week": 5,
    "month": 12
  }'

Request URL
http://127.0.0.1:8000/predict

Server response
Code Details
200 Response body
{
  "district": 11,
  "risk_score": 0.5598,
  "alert": true,
  "action": "DISPATCH"
}
Download

Response headers
Content-length: 68
Content-type: application/json
Date: Tue, 09 Dec 2025 22:39:45 GMT
Server: uvicorn

Responses
Code Description Links
200 Successful Response No links
Media type
application/json
Controls Accept header.

```

## 2. Technical Documentation

### System Architecture

The project follows a modular **MLOps (Machine Learning Operations)** architecture, separating training, serving, and monitoring into distinct components.

- **Model Type:** Histogram-based Gradient Boosting Classifier (`HistGradientBoostingClassifier`).
- **Frameworks:** Scikit-Learn (Modeling), FastAPI (Serving), Pandas/Numpy (Data Processing).
- **Input Data:** Historical crime records (Date, Primary Type, District).
- **Target Variable:** Binary Classification (`Is_Violent`: 1 for Battery/Assault/Robbery/Homicide, 0 for others).

## Data Pipeline & Feature Engineering

Raw data is transformed before being fed into the model to capture temporal and spatial patterns:

Feature	Transformation Logic	Purpose
District	Categorical Encoding	Captures spatial risk profiles of different neighborhoods.
Hour	Sine/Cosine Transformation	Converts linear time (0-23) into cyclical features to model day/night cycles.
Month	Sine/Cosine Transformation	Captures seasonal trends (e.g., crime spikes in summer).
Day of Week	Integer (0-6)	Captures weekend vs. weekday patterns.

## Logging & Monitoring Strategy

To ensure the model remains accurate over time, the system implements a continuous logging mechanism:

- **Live Logs** ([live\\_logs.csv](#)): Every request sent to the API is automatically recorded in this CSV file. It captures the input features (District, Time) and the model's output (Risk Score).

## Model Performance Strategy

Due to severe class imbalance (violent crimes are rarer than non-violent ones), standard accuracy was a misleading metric. The model was optimized using:

1. **Balanced Class Weights:** Increasing the penalty for missing a violent crime during training.
2. **Recall Optimization:** Hyperparameters were tuned to maximize Recall (catch rate) rather than Precision.
3. **Threshold Moving:** The decision boundary was lowered from the default 50% to an optimal value to ensure >65% of violent crimes are flagged.

## Final Metrics:

- **Recall (Violence):** ~65% (vs. 3% baseline)
- **ROC-AUC:** ~0.62

- **Drift Metric:** Jensen-Shannon Distance.
  - **F1\_SCORE:** 0.4581
- 

## 3. Deployment Instructions

### Prerequisites

- Python 3.9+ installed.
- Required libraries: `pandas`, `scikit-learn`, `fastapi`, `unicorn`, `joblib`.

### Installation Steps

#### 1. Clone the Repository & Setup Environment

Create a project folder and install dependencies:

Bash

```
mkdir police_project  
cd police_project  
pip install pandas scikit-learn fastapi unicorn joblib scipy
```

#### 2. Train the Model

Run the training pipeline to generate the model artifact (`police_model_auto.pkl`):

Bash

```
python police_ops_auto.py
```

- *Output:* A serialized model file is created in the directory.

#### 3. Launch the API Server

Start the FastAPI server using Unicorn:

Bash

```
unicorn main:app --reload
```

- *Success Message:* `INFO: Unicorn running on http://127.0.0.1:8000`

#### 4. Verify Deployment

Open a web browser and visit `http://127.0.0.1:8000/docs`. If the dashboard loads, the deployment is successful.

## Production Considerations

- **Data Persistence:** Ensure the `live_logs.csv` file is backed up regularly, as it contains the history of all predictions made by the system.